

WireGuard-go over VSOCK

Guilherme Balena Versiani

<guilherme.versiani@zerohash.com>

ZeroHash LLC

April 3, 2024

Abstract

WireGuard is a modern VPN protocol that was originally written in C. Due to the increasing popularity of the Go programming language, the need for a Go implementation of the protocol arose. In early 2018, WireGuard’s creator Jason Donenfeld announced on the WireGuard mailing list that he was working on a new Go implementation of the protocol called WireGuard-go. The goal of WireGuard-go was to provide a lightweight and efficient implementation of the WireGuard protocol that could be easily integrated into various projects and platforms.

The purpose of this paper is to provide an overview of the WireGuard-go usage to provide AWS Nitro Enclave’s networking over an AWS Elastic Kubernetes Service (EKS) through VSOCK.

and secure communication between virtual machines (guests) and the host [12]. It was designed to overcome the limitations of traditional networking protocols when it comes to performance, latency, and scalability in virtualized environments. It provides a high-performance communication channel that is optimized for virtualized environments and is designed to avoid the overhead associated with traditional network protocols. It was first introduced in the Linux kernel in 2013 by VMware, and since then it has been used in their VMware ESXi hypervisors, VMware Fusion and VMware Workstation products. It is also supported in some cloud-based VMware offerings, such as VMware Cloud on AWS and VMware Cloud Director. QEMU[3], a popular open-source virtualization platform, and OpenShift, a container platform that includes virtualization capabilities based on KubeVirt [10], are other examples of systems using VSOCK.

1 Background

AWS Nitro Enclaves provide a highly secure and isolated environment for running sensitive workloads. However, due to their isolation, Nitro Enclaves do not have network access by default [2]. This can pose a challenge for applications that require network connectivity, such as those running in Kubernetes clusters. To address this challenge, we designed a WireGuard-go extension that supports VSOCK to provide full network access to Nitro Enclaves.

1.1 VSOCK

VSOCK, also known as vSocket or VMware Socket Direct Path I/O, is a socket-based communication mechanism that enables efficient

The `vsock` module was integrated in the Linux Kernel 3.9 [15], so it is typically available on most modern Linux distributions. When loaded, applications can create `AF_VSOCK` sockets to communicate [8]. Each VM is assigned a unique CID (Context ID) that is used to identify it on the vsock network, while the host has a unique identifier (CID 2). Valid socket types are `SOCK_STREAM` and `SOCK_DGRAM`. `SOCK_STREAM` provides connection-oriented byte streams with guaranteed, in-order delivery. `SOCK_DGRAM` provides a connectionless datagram packet service with best-effort delivery and best-effort ordering. Availability of these socket types is dependent on the underlying hypervisor. AWS Nitro Enclaves hypervisor in particular does not support `SOCK_DGRAM` socket

type. Sockets are affected by live migration of virtual machines. Connected `SOCK_STREAM` sockets become disconnected when the virtual machine migrates to a new host. Applications must reconnect when this happens.

1.2 WireGuard

Among the most popular Open Source VPN software, WireGuard is known for its high performance and low resource usage. It has a smaller codebase compared to other VPN solutions, which makes it faster and easier to audit. WireGuard is also designed to use modern cryptography that is more efficient, resulting in lower CPU usage and longer battery life on mobile devices. As a comparative, the popular OpenVPN is generally considered to have a higher resource usage compared to WireGuard. It requires more CPU cycles and memory to establish and maintain VPN connections due to its more complex protocol and encryption algorithms. StrongSwan, another popular Open Source VPN software, has also higher resource usage compared to both WireGuard and OpenVPN [5]. As of March 2023, there have been 1 CVE attributed to WireGuard (for Windows), 56 CVEs attributed to OpenVPN, and 47 CVEs attributed to strongSwan. WireGuard uses a variety of cryptographic primitives, including the Noise Protocol Framework, which provides a secure way to establish a shared secret key between the two devices, and a combination of modern encryption algorithms, such as ChaCha20, Poly1305, BLAKE2s, and Curve25519.

Our approach leverages WireGuard-go [7], a Go implementation of WireGuard, to enable secure connectivity between the Nitro Enclave and Kubernetes. WireGuard-go was created to provide a high-performance, portable, and easy-to-use implementation of the WireGuard protocol in the Go programming language, with the goal of making it easier for developers to create and deploy secure VPN solutions on a wide range of platforms and devices.

Unlike other VPN protocols, WireGuard explicitly supports UDP only [6], which enables it to operate efficiently and avoid the problems

associated with TCP. However, when interfacing with VSOCK on AWS Nitro Enclaves, there is no support for `SOCK_DGRAM`. Because of that, the WireGuard-go-vsock [11] code creates an optimized `SOCK_STREAM` channel between the VSOCK peers, and over that, normal WireGuard datagrams are exchanged. It is also worth mentioning that due to address differences, the CID formed by a 32-bit integer gets mapped to an IPv4 address at the WireGuard-go stack, and the port, that is also another 32-bit integer, gets mapped into a 16-bit port to match TCP/IP ports.

It is important to note that a WireGuard VPN between a Kubernetes container and an Enclave does not necessarily mean that the Enclave has access to the local network, as WireGuard uses TUN devices [13]. Instead, it relies on IP packet forwarding based on routing rules and Network Address Translation (NAT). The WireGuard VPN creates a secure tunnel between the Enclave and the Kubernetes container only, and traffic between the Enclave and the outside world is still subject to routing rules and other network configurations.

Establishing a VPN through VSOCK between a host and an AWS Nitro Enclave allows for direct network communication between the host and the Enclave through “cryptographic tunnels”, where all communication is encrypted and transmitted through a virtual network interface (like `wg0`) transparently. This interface can be seen as a secure point-to-point link between two devices, and all traffic is directed to it. This setup does not require a proxy, which is usually the approach offered by the AWS SDK. The Amazon’s `vsock-proxy` requires adding a traffic forwarder at the Enclave side, alongside with the Enclave app, and configuring client libraries to use the local traffic forwarder [1]. The other drawback of using `vsock-proxy` is that it requires launching one instance per remote address of the server to be proxied. The VPN simplifies the setup of libraries, as there is no need to configure an HTTP proxy in each, and management of the system, because it doesn’t require multiple proxies one per service, while still maintain security by keeping the attack surface small constrained to the WireGuard peers

and networking set at both sides.

Another project adopting a similar approach is `nitriding` [14], with the exception they don't implement a VPN. But the outcome is similar, as they claim:

Nitriding creates a TAP interface inside the enclave, allowing your application to transparently access the Internet without having to worry about VSOCK, port forwarding, or tunneling.

In this approach, a peer-to-peer network link is established between the Enclave and a controlling application, known as the Relay, rather than relying on a traditional proxy setup. The Relay conveys container configurations — including IP address, DNS search parameters, and default gateway settings — back to the Enclave application, all while facilitating the establishment of a peer-to-peer VPN via VSOCK. Unlike `nitriding`, the Relay possesses the added capability of attesting the Enclave application's integrity before permitting traffic to flow to or from it. This additional layer of security significantly enhances the solution's robustness, rendering Enclave tampering far less likely to occur.

1.3 Kubernetes

The Relay is hosted by an AWS EKS (Kubernetes) environment, and deployed using regular manifests, except they contain a few Nitro Enclave extensions. These extensions indicates to Kubernetes where the Pods should be allocated, case there are managed EKS nodes without Enclaves support and self-managed ones in the same cluster, and also enable the `/dev/nitro_enclaves` device on the running containers, so that they can launch/terminate Enclaves. In order to be able to manage NAT and port forwarding for the Enclave, the Relay container also has to have the `NET_ADMIN` capability.

One of the concerns that arise from the `NET_ADMIN` enablement is that firewall rules would conflict with `kube-proxy` or CNI. But

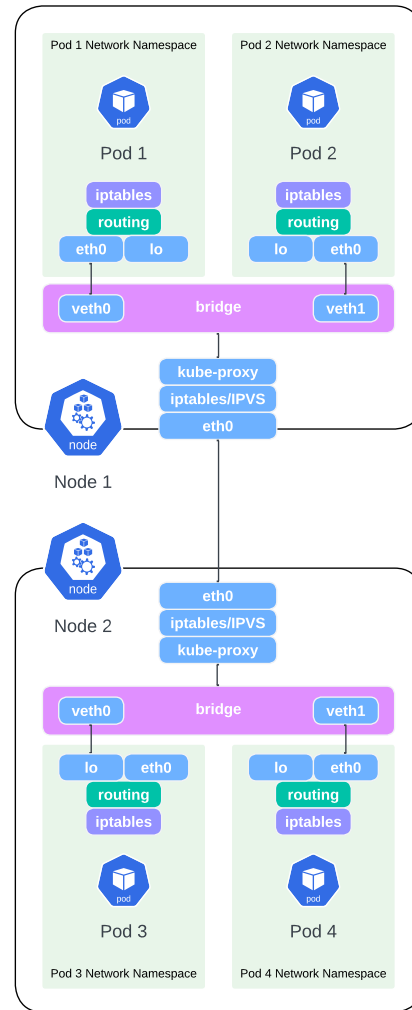


Figure 1: Kubernetes networking.

this is not the case. Kubernetes provides network isolation to Pods through various mechanisms implemented in the Linux kernel itself. Each Pod is assigned a unique IP address within the Kubernetes cluster, and this is achieved through the use of *network namespaces* [4]. Network namespaces provide a way to isolate network resources such as interfaces, routes, and firewall rules, as shown in figure 1.

When a container is launched within a Kubernetes Pod, the Kubernetes CNI (Container Network Interface) plugin creates a new network namespace for the Pod. A *new* virtual network device called `lo` (loopback device), which is similar to the one in the default network namespace except it is isolated, is set at the namespace. It also creates a new pair of network devices, a virtual Ethernet device called `veth` residing in the parent namespace, and the other end residing in the new child network namespace, usually `eth0`. In order for them to work, the CNI connects the `veth` device to a bridge or a physical network interface in the parent namespace, allowing the container to communicate with the outside world. In addition, it can also use Linux's firewall (`iptables` or `bpfilter`) to implement network policies at the Node level for the Pod. These rules are invisible to the Pods except when they get `hostNetwork` access.

Kubernetes also supports various network plugins that allow for different network topologies and configurations, such as Cilium. These plugins are responsible for implementing the network model of the cluster, and they can be used to provide additional network isolation and security. Due to the usage of network namespaces, none of these plugins conflict with the `iptables` rules used by the Relay.

1.4 Security

The `NET_ADMIN` capability allows a container to perform various networking-related tasks that are otherwise restricted in a default container configuration. Some of the threats associated with a container configured with `NET_ADMIN` capability include:

- **Network sniffing and eavesdropping.** A container with `NET_ADMIN` capability can

be used to sniff network traffic passing through the host system's network interfaces. Attackers can use this capability to capture sensitive information, such as credentials and other sensitive data, as it traverses the network. By design, Enclave's data in transit is encrypted, which mitigates this risk for Enclave's communication, but other services not properly secured can be compromised.

- **Network injection and packet spoofing.** With `NET_ADMIN` capability, a container can inject or modify network packets as they traverse the network. Attackers can use this capability to launch various network attacks, such as man-in-the-middle attacks or packet spoofing attacks.
- **Network enumeration and reconnaissance.** A container with `NET_ADMIN` capability can be used to scan the host system's network interfaces and map out its topology. Attackers can use this capability to identify potential targets and vulnerabilities within the network.
- **Network DoS attacks.** With `NET_ADMIN` capability, a container can be used to launch various types of network-based denial-of-service (DoS) attacks against the host system or other network resources.

To configure WireGuard without granting `NET_ADMIN` capabilities, one would use gVisor's networking stack to handle the WireGuard setup within the container. gVisor can intercept the necessary system calls made by the container to set up WireGuard interfaces, manage routes, and handle firewall rules. This way, the container can perform WireGuard configuration tasks securely within its isolated environment without requiring escalated privileges.

On the other hand, the benefits may outweigh the risks:

- Enclave's data in transit is encrypted, so network sniffing risks are mitigated.

- Enclaves are attested before use, communication is immune to network injection and packet spoofing.
- Because of the above, network enumeration and reconnaissance doesn't expose sensitive information about internal systems; at the same time Enclave's network devices are isolated and cannot be enumerated from the Node.
- As the network in question is private, DoS attacks are less likely to be implemented without being rapidly identified.

In any case, it is recommended to implement other security measures such as network segmentation and access control because non Enclave services don't offer any of the isolation characteristics of the Enclave.

2 Architecture

The AWS EKS cluster is created in a VPC. By default, Pod networking is provided by the Amazon VPC Container Network Interface (CNI) plugin, which can be replaced by Cillium. Every CNI plugin allows Kubernetes Pods to have the same IP address as they do on the VPC network. In this regard, there is nothing special at Cillium when compared to Amazon CNI.

Linux kernel network namespaces are used to isolate network resources between containers and Pods. When creating a Pod, several new namespaces are used for isolation, including a network namespace to isolate the network of the container. The network namespace provides a virtual network stack to the container, which includes its own network interfaces, routing tables, and `iptables` rules. This allows containers to have their own network namespace, which is separate from the host network namespace and other containers.

Having the containers network isolated, the WireGuard peer implemented at the Relay can make use of `iptables` to masquerade traffic from the Enclave and ingress traffic to the Enclave by performing a destination NAT on the incoming packets.

The syntax for these rules are presented in listings 1 and 2.

Listing 1 Enclave Egress

```
1 iptables -t nat -A POSTROUTING \
2   -s 100.64.254.0/24 \
3   -o eth0 \
4   -j MASQUERADE
```

Listing 2 Enclave Ingress

```
1 iptables -t nat -A PREROUTING \
2   -i eth0 -p tcp --dport 8080 \
3   -j DNAT \
4   --to-destination 100.64.254.2
```

The Enclave-Relay architecture is presented in figure 2. During initialization, the Relay container uses `nitro-cli` to launch the Enclave prior to starting its own Relay process. The Relay and the Enclave applications start a WireGuard-go stack without any peer using random key pairs. Then an auto provisioning step is performed at the next VSOCK port of the WireGuard's one (by default the WireGuard port is 51820 so the auto provisioning happens at the 51821), where the Relay acts as an initiator, and the Enclave as a responder. The auto provisioning initiator connects to the responder via VSOCK and offers all local network settings, including its WireGuard-go public key and a `nonce` that is used at the attestation document as a proof of authenticity of the Enclave. The Enclave configures its local network, the loopback interface `lo`, the local IP of the `wg0` interface, the default gateway is set to the Relay IP, and the DNS settings are set as a copy from the Relay's one. The Enclave WireGuard-go stack gets set with a peer entry corresponding to the Relay public key along with other settings. As a response to the auto provisioning request, the Enclave sends an attestation document containing the same `nonce` of the request, and it's WireGuard-go's public key. The initiator receives the attestation document and validates it; if valid, it sets up the Enclave's peer at the WireGuard-go stack and the usual WireGuard protocol handshake starts, assigning a shared key at both ends of the connection.

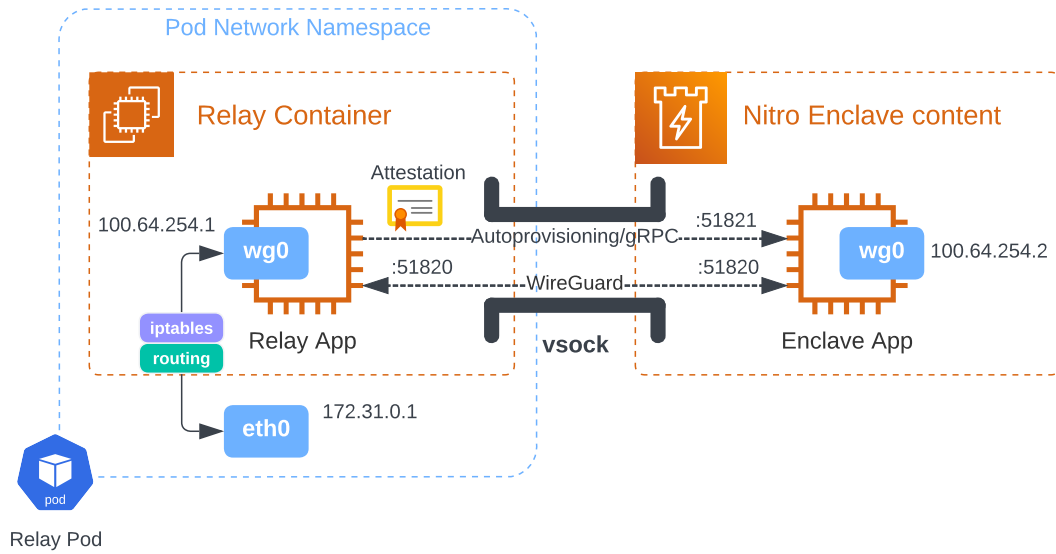


Figure 2: Enclave-Relay architecture.

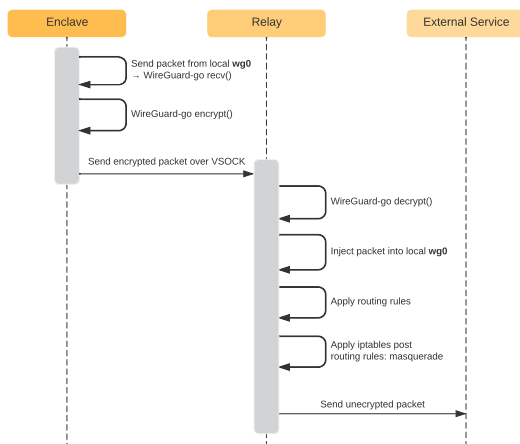


Figure 3: Enclave Egress sequence.

2.1 Enclave Egress Traffic

The packet lifecycle from the perspective of the Enclave when it needs to communicate with an external service is shown at the figure 3, and described below:

1. The packet originates from the Enclave application hosted in an AWS Nitro Enclave using the address assigned to `wg0` as the

source address according to the local routing rules, as the default gateway set at the Enclave is the Relay address set at the WireGuard peer-to-peer connection. The prepared WireGuard-go code is listening for packets on that `wg0` TUN interface, so it intercepts the packet.

2. The WireGuard-go code uses the shared key established after the protocol handshake to encrypt the packet.
3. The encrypted packet is sent over VSOCK, as initially provisioned by the Relay.
4. The Relay application receives the encrypted packet from the VSOCK interface, and delivers it to WireGuard-go, which decrypts the packet using the shared key established during the protocol handshake.
5. The packet is injected into the local `wg0` interface, as if it would have been received from the Enclave through a regular network interface.
6. The local routing rules set at the Pod network namespace will apply, meaning that

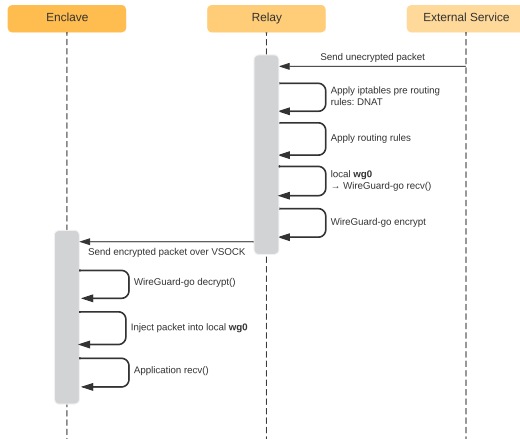


Figure 4: Enclave Ingress sequence.

the local default gateway will be used to dispatch the packet.

7. Then the local `iptables` rules also set at the Pod network namespace only will apply. As it has been set to masquerade packets from the Enclave, the source address switches to the Relay's local IP set at the `eth0`.
8. The Linux Kernel now sends the packet through the `eth0` network namespace of the Pod, which is set by the Kubernetes CNI to bridge, VPN or directly connect to the Node's `eth0` interface, reaching out to the external service.

2.2 Enclave Ingress Traffic

The packet lifecycle from the perspective of the Enclave when it receives data from an external service is showed at the figure 4, as described below:

1. According to the Kubernetes CNI, the packet is sent to the local `eth0` of the Relay's Pod.
2. The `iptables` DNAT rules set at the Pod network namespace are applied, so the destination address is set to the IP of the Enclave behind the VSOCK.

3. The routing rules set at the Pod network namespace are also applied, which indicates that the interface to be used to dispatch the packet is the local `wg0` interface.
4. The local TUN `wg0` interface is attached to the Relay's WireGuard-go.
5. At the application level, WireGuard-go encrypts the packet with the shared key exchanged during the protocol handshake.
6. The Relay sends the encrypted packet to the Enclave over VSOCK.
7. The WireGuard-go stack set at the Enclave side receives the packet and decrypts it with the shared key exchanged during the protocol handshake.
8. The WireGuard-go then injects the unencrypted packet into the local `wg0` TUN interface, as if it were just received from a regular network interface.
9. As the destination address of the packet is the local IP set at the `wg0`, the Linux Kernel then delivers the packet to the Enclave application.

3 Conclusion

This paper presented the solution for using WireGuard-go over VSOCK in an Elastic Kubernetes Service (EKS) with support to AWS Nitro Enclaves. We demonstrated that Kubernetes already provides network isolation so that `iptables` commands do not interfere with it. The solution securely connects Enclave-based applications with external networks using WireGuard encrypted traffic. We showed how the implementation of WireGuard-go over VSOCK provides a secure and efficient communication channel for Enclave-based applications.

References

- [1] AMAZON, I. Vsock proxy. <https://github.com/aws/aws-nitro-enclaves-cli/blob/>

- main/vsock_proxy/README.md, 2023. Accessed: March 20, 2023.
- [2] AMAZON, I. What is aws nitro enclaves? <https://docs.aws.amazon.com/enclaves/latest/user/nitro-enclave.html>, 2023. Accessed: March 20, 2023.
- [3] BELLARD, F. Qemu features/virtioVsock. <https://wiki.qemu.org/Features/VirtioVsock>, 2023. Accessed: March 20, 2023.
- [4] CRITELLI, A. Kubernetes network stack fundamentals: How containers inside a pod communicate. <https://www.redhat.com/sysadmin/kubernetes-pod-network-communication>, 2022. Accessed: March 20, 2023.
- [5] DEKKER, E., AND SPAANS, P. A performance comparison of the vpn implementations wireguard, strongswan and openvpn in a one gbit/s environment. <https://rp.os3.nl/2019-2020/p71/presentation.pdf>, 2019-2020. Accessed: March 20, 2023.
- [6] DONENFELD, J. Known limitations - wireguard. <https://www.wireguard.com/known-limitations/>, 2015. Accessed: March 20, 2023.
- [7] DONENFELD, J. Go implementation of wireguard. <https://git.zx2c4.com/wireguard-go>, 2023. Accessed: March 20, 2023.
- [8] KERRISK, M. vsock(7) - linux manual page. <https://man7.org/linux/man-pages/man7/vsock.7.html>, 2021. Accessed: March 20, 2023.
- [9] LLC, G. gvisor, application kernel for containers. <https://github.com/google/gvisor>. Accessed: April 3, 2024.
- [10] RED HAT, I. Red hat openshift sandboxed containers: Peer-pods technical deep dive. https://kubevirt.io/user-guide/virtual_machines/vsock/, 2023. Accessed: March 20, 2023.
- [11] VERSIANI, G. Go implementation of wireguard over vsock. <https://github.com/seedcx/wireguard-go-vsock>, 2023. Accessed: March 20, 2023.
- [12] VMWARE, I. vsockets programming guide. https://vdc-download.vmware.com/vmwb-repository/dcr-public/a49be05e-fa6d-4da1-9186-922fbfef149e/a65f3c51-aaeb-476d-80c3-827b805c2f9e/ws9_esx60_vmci_sockets.pdf, 2023. Accessed: March 20, 2023.
- [13] WIKIPEDIA, T. F. E. Tun/tap. <https://en.wikipedia.org/wiki/TUN/TAP>, 2023. Accessed: March 20, 2023.
- [14] WINTER, P., GILES, R., DAVIDSON, A., AND PESTANA, G. A framework for building secure, scalable, networked enclaves, 2022.
- [15] ZHANG, G. Vsock for linux upstreaming. <https://lwn.net/Articles/520000/>, 2012. Accessed: March 20, 2023.